

reinforcement


October 20, 2025

1 Reinforcement Learning

1.1 By Noël Jumin and Louis Maurin

1.1.1 October 9, 2025

```
[1]: import numpy as np
```



We define the transition matrices as follows:

$$T(S, a0, S') = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 0 & 1-x & 0 & x \\ 1-y & 0 & 0 & y \\ 1 & 0 & 0 & 0 \end{bmatrix}$$

$$T(S, a1, S') = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

$$T(S, a2, S') = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

With the reward vector:

$$R = [0, 0, 1, 10]$$

2 Q1

We can so write the possible policies as:

$$\pi_1 : \begin{cases} S_0 \rightarrow a_1 \\ S_1 \rightarrow a_0 \\ S_2 \rightarrow a_0 \\ S_3 \rightarrow a_0 \end{cases} \quad \pi_2 : \begin{cases} S_0 \rightarrow a_2 \\ S_1 \rightarrow a_0 \\ S_2 \rightarrow a_0 \\ S_3 \rightarrow a_0 \end{cases}$$

3 Q2

We can write the Bellman equations for the optimal value function (V^*) as follows:

$$V^*(S) = R(S) + \gamma \max_a \sum_{S'} T(S, a, S') V^*(S')$$

Where (γ) is the discount factor.

To solve for (V^*), we can set up the equations for each state:

1. For state 0:

$$\begin{aligned} V^*(S_0) &= R(S_0) + \gamma \max(V^*(S_1), V^*(S_2)) \\ V^*(S_0) &= \gamma \max(V^*(S_1), V^*(S_2)) \end{aligned}$$

2. For state 1:

$$\begin{aligned} V^*(S_1) &= R(S_1) + \gamma((1-x)V^*(S_1) + xV^*(S_3)) \\ V^*(S_1) &= \gamma((1-x)V^*(S_1) + xV^*(S_3)) \end{aligned}$$

3. For state 2:

$$\begin{aligned} V^*(S_2) &= R(S_2) + \gamma \max((1-y)V^*(S_0) + yV^*(S_3)) \\ V^*(S_2) &= 1 + \gamma((1-y)V^*(S_0) + yV^*(S_3)) \end{aligned}$$

4. For state 3:

$$V^*(S_3) = R(S_3) + \gamma \max(V^*(S_0)) V^*(S_3) = 10 + \gamma V^*(S_0)$$

4 Q3

Now that we have these equations, we could determine a value for x , that for all $\gamma \in [0, 1[$, and $y \in [0, 1]$, $\pi^*(S_0) = a_2$.

We know that

$$\pi^*(S_0) = \arg \max_a \sum_{S'} T(S_0, a, S') V^*(S')$$

Hence,

$$\pi^*(S_0) = \begin{cases} a_1 & \text{if } V^*(S_1) > V^*(S_2) \\ a_2 & \text{if } V^*(S_1) < V^*(S_2) \end{cases}$$

So we need to find x so that $V^*(S_2) > V^*(S_1)$

To do that, we'll take $x = 0$ and reformulate $V^*(S_1)$ and $V^*(S_2)$ Bellman equations (Taking into account that $R(S)$ is not function of a) :

$$\begin{aligned} V^*(S_1) &= \gamma((1-x)V^*(S_1) + xV^*(S_3)) \\ V^*(S_1) &= \gamma V^*(S_1) \\ (1-\gamma)V^*(S_1) &= 0 \\ V^*(S_1) &= 0 \quad \text{since } \gamma \in [0, 1[\end{aligned}$$

$$\begin{aligned} V^*(S_2) &= \gamma((1-y)V^*(S_0) + yV^*(S_3)) \\ V^*(S_2) &= \gamma((1-y)V^*(S_0) + y\gamma V^*(S_0)) \\ V^*(S_2) &= \gamma V^*(S_0)(1-y+y\gamma) \\ V^*(S_2) &= \gamma V^*(S_0)(1+y(\gamma-1)) \end{aligned}$$

So $V^*(S_2) > V^*(S_1)$ for any $y \neq 0$ and $x = 0$ is a solution.

5 Q4

We can now try to determine a value for y , that for all $\gamma \in [0, 1[$, and $x > 0$, $\pi^*(S_0) = a_1$.

As of before,

$$\pi^*(S_0) = \begin{cases} a_1 & \text{if } V^*(S_1) > V^*(S_2) \\ a_2 & \text{if } V^*(S_1) < V^*(S_2) \end{cases}$$

Let's suppose such a value of y exists. We call it y_0 and choose $y = y_0$.

We now have that for all $\gamma \in [0, 1[$, and $x > 0$, $\pi^*(S_0) = a_1$ which is equivalent to $V^*(S_1) > V^*(S_2)$.

We can that write,

$$\begin{aligned} V^*(S_0) &= \gamma V^*(S_1) \\ V^*(S_3) &= 10 + \gamma^2 V^*(S_1) \end{aligned}$$

So,

$$\begin{aligned}
 V^*(S_1) &= \gamma(1-x)V^*(S_1) + \gamma x(10 + \gamma^2 V^*(S_1)) \\
 V^*(S_1)(1 + \gamma(x-1) - x\gamma^3) &= 10x\gamma \\
 V^*(S_1) &= \frac{10x\gamma}{1 + \gamma(x-1) - x\gamma^3} \quad (\text{for } 1 + \gamma(x-1) - x\gamma^3 \neq 0)
 \end{aligned}$$

Now using this formula, if $\gamma = 0$ then $1 + \gamma(x-1) - x\gamma^3 = 1 \neq 0$ and $V^*(S_1) = 0$.

Actually, we already had from the equations above that. When $\gamma = 0$,

$$\begin{aligned}
 V^*(S_1) &= 0 \\
 V^*(S_2) &= 1
 \end{aligned}$$

So $V^*(S_2) > V^*(S_1)$ and we end up with a contradiction.

So such a value of y doesn't exist.

6 Q5

We'll now implement our value iteration:

```
[6]: R = [0, 0, 1, 10]
gamma = 0.9
x = 0.25
y = 0.25
tol = 1e-6

# Transition matrices
T_a0 = np.array([
    [0.0, 0.0, 0.0, 0.0],
    [0.0, 1.0 - x, 0.0, x],
    [1.0 - y, 0.0, 0.0, y],
    [1.0, 0.0, 0.0, 0.0],
])
T_a1 = np.array([
    [0.0, 1.0, 0.0, 0.0],
    [0.0, 0.0, 0.0, 0.0],
    [0.0, 0.0, 0.0, 0.0],
    [0.0, 0.0, 0.0, 0.0],
])
T_a2 = np.array([
    [0.0, 0.0, 1.0, 0.0],
    [0.0, 0.0, 0.0, 0.0],
    [0.0, 0.0, 0.0, 0.0],
    [0.0, 0.0, 0.0, 0.0],
])
T = [T_a0, T_a1, T_a2]
```

```

n_states = 4 # S0, S1, S2, S3
n_actions = len(T) # One action per transition matrix

V = np.zeros(n_states)
delta = np.inf
while delta > tol:
    V_prev = V.copy()
    for s in range(n_states):
        q_vals = [np.dot(T[a][s], V_prev) for a in range(n_actions)]
        V[s] = R[s] + gamma * max(q_vals)
    delta = np.max(np.abs(V - V_prev))

# Extract deterministic optimal policy
action_names = ['a0', 'a1', 'a2']
pi = []
for s in range(n_states):
    q_vals = [np.dot(T[a][s], V) for a in range(n_actions)]
    print(f" S{s}: q = {np.round(q_vals,6)} -> best = {action_names[int(np.
    ↪argmax(q_vals))]}")
    pi.append(action_names[int(np.argmax(q_vals))])

print("V* =", np.round(V, 6))
print("pi* =", pi)

```

```

S0: q = [ 0.          15.761813 15.69789 ] -> best = a1
S1: q = [17.513126  0.          0.          ] -> best = a0
S2: q = [16.33099  0.          0.          ] -> best = a0
S3: q = [14.185631  0.          0.          ] -> best = a0
V* = [14.185631 15.761813 15.69789 22.767067]
pi* = ['a1', 'a0', 'a0', 'a0']

```

To explain what appened here, V^* is the calculation of reward taking in account if it will be immediate reward or not immediate but lead to a big reward.

With Q , we can know from each state, what action take to get the maximum reward and so take the best path.

With this information we know that the best option from S_0 is a1 and a0 for the others.